

# Cluster Failure: fMRI's Big Shake-Up

Chris Hammill

September 13, 2016

Tens Of Thousands Of fMRI Brain Studies May Be Flawed



# Bug in fMRI software calls 15 years of research into question

Popular pieces of software for fMRI were found to have false positive rates up to 70%

Science News

from research.org.uk

## Software for fMRI yield erroneous results

Cluster failure: Why fMRI inferences for spatial extent have inflated false positive rates

OOPSIE! —

## Software faults raise questions about the validity of brain studies

Interpretation of functional MRI data called into question.

JOHN TIMMER - 7/1/2016, 2:55 PM

# Cluster failure: Why fMRI inferences for spatial extent have inflated false-positive rates

Anders Eklund<sup>a,b,c,1</sup>, Thomas E. Nichols<sup>d,e</sup>, and Hans Knutsson<sup>a,c</sup>




# So What Happened

- ▶ Eklund, Nichols, and Knutsson demonstrated standard fMRI statistical inference has badly inflated false positives rates
- ▶ Makes you wonder if exciting brain region X responding to stimulus Y finding was just a cherry-picked false positive.
- ▶ Highlighted that due to non-reproducible workflows, and poor data sharing, many of these finding could never be repeated with valid inference.

# How Did We Get Here

- ▶ fMRI is challenging to analyze
- ▶ Preprocessing steps widely used as black boxes
- ▶ Desire to use spatial information to determine signal significance
- ▶ Improperly specified models of spatial noise
- ▶ Ultimately a multiple testing problem

# Multiple Comparisons



# Multiple Comparisons

- ▶ As with most imaging analysis, multiple comparisons is significant concern
- ▶ Solutions:
  1. Bonferroni: control your type one error rate by multiplying your results by the number of tests. This is equivalent to setting your type one error rate to  $\alpha/n$
  2. FDR (Benjamini-Hochberg): Order your p-values lowest to highest and set  $p_{corr} = \frac{pn}{rank(p)}$
  3. FDR (Variants): Assume p-values come from two distributions, a true finding distribution peaked near zero, and a uniform distribution.
- ▶ But in low power situations with high covariance between voxels, these can be conservative

# Correction To The Paper

*The authors note that on page 7900, in the Significance Statement, lines 9–11, “These results question the validity of some **40,000** fMRI studies and may have a large impact on the interpretation of neuroimaging results” should instead appear as “These results question the validity of a number of fMRI studies and may have a large impact on the interpretation of weakly significant neuroimaging results.”*

Nichols goes on to say on his blog, the number is closer to 15,000 with an additional 13,000 papers using *no multiple comparison correction at all*

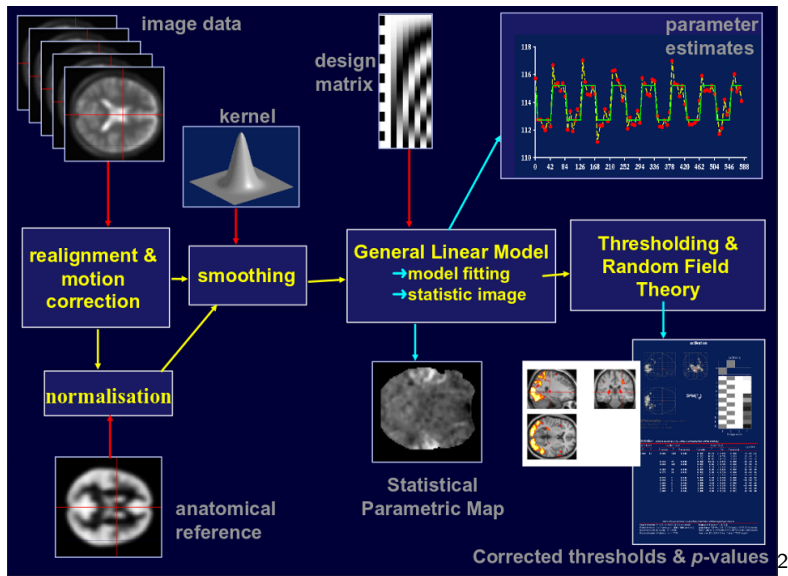


# About Group Comparisons Task-Based fMRI

- ▶ Most fMRI seeks to measure brain activity by blood flow
- ▶ Blood oxygen level dependent (BOLD) contrast
- ▶ A time-series of volumes are acquired for each subject
- ▶ Stimuli are presented to the subject throughout the time series
- ▶ The BOLD signal is modelled as a function of the stimuli
- ▶ The statistical associations of the BOLD contrast to the stimuli are compared across groups
- ▶ Group comparisons typically simple t-tests and ANOVAs


# Why Is This Tough

- ▶ Subjects move:
  - ▶ within subject each fMRI volume must be aligned to each-other
  - ▶ these must be aligned to a corresponding anatomical scan
  - ▶ these must be registered to a common space
- ▶ BOLD signal is sluggish
  - ▶ ~ 2 seconds to start
  - ▶ ~ 4-6 to peak
  - ▶ ~ 10 to return to baseline so the stimulus time series is convolved with a function to match this behaviour
- ▶ Analyzing time series comes with it's own statistical challenges
  - ▶ how do we model temporal autocorrelation



# Enter Spatial Models

- ▶ Signals with large spatial extent are probably more likely to be real than individual high intensity
- ▶ Question becomes, how do we analyze spatial extent, and how do we correct for multiple comparisons?
- ▶ First: Threshold your data at threshold that sounds appealing ( $p < .01$ , and  $p < .001$ )
- ▶ Then: Use random field theory (RFT) results to assign a p-value to clusters based on their size





# The Problems

- ▶ When statistics maps aren't smooth enough, RFT p-values are biased (2003)
- ▶ RFT typically assumes a stationary noise distribution (same noise over the brain) which is often invalid (2004)
- ▶ Together these problems can lead to 70% FWE rates in single subject analyses (2012)



# The Paper

- ▶ In order to assess how much these problems matter for group comparisons, check the null distribution
- ▶ The authors took a large open data set with a pool of neurotypical subjects, and randomly sampled groups to compare
- ▶ If after processing and multiple comparison correction any clusters in the brain were significant that test was a false positive (error).
- ▶ The distribution for a two group difference should be Student's t distribution, and after bonferroni correction, the expected proportion of errors should be 5%
- ▶ Higher error rates imply the multiple comparison correction is insufficient.
- ▶ Five analysis functions from the three most popular fMRI software packages were compared to their non-parametric alternative

A





B





# The results

- ▶ All parametric tools produce FWE higher than 5%
- ▶ Situation is more extreme when cluster defining thresholds are high (FWE rates ~20-40)
- ▶ Different data sets are affected differently (Beijing less affected than Cambridge)



# So What To Do

- ▶ Give up on cluster inference
- ▶ Use a bootstrap/randomization test instead of RFT
  1. Shuffle group membership and covariates between subjects
  2. Refit your model
  3. Create a distribution on largest clusters per sample
  4. Assign p-values to cluster size from this null distribution
- ▶ Merge voxel level data with cluster extent (Threshold free cluster enhancement)
  1. Shuffle group membership and covariates between subjects
  2. Refit your model
  3. Perform TFCE
  4. Create a randomization of the cluster enhanced statistics
  5. Assign p-values to voxels from this null distribution
  6. Correct with Bonferroni or FDR

## Questions and Discussion